

Project I - Convolutional neural networks

Deep Learning 2024

Kinga Frańczak, 313335

Grzegorz Zakrzewski, 313555

Contents

1. Introduction

2. Experiments

3. Results

4. Conclusions

Description of the research problem



Figure: Sample images from each class in the CINIC-10 dataset.

Experiments - network architectures

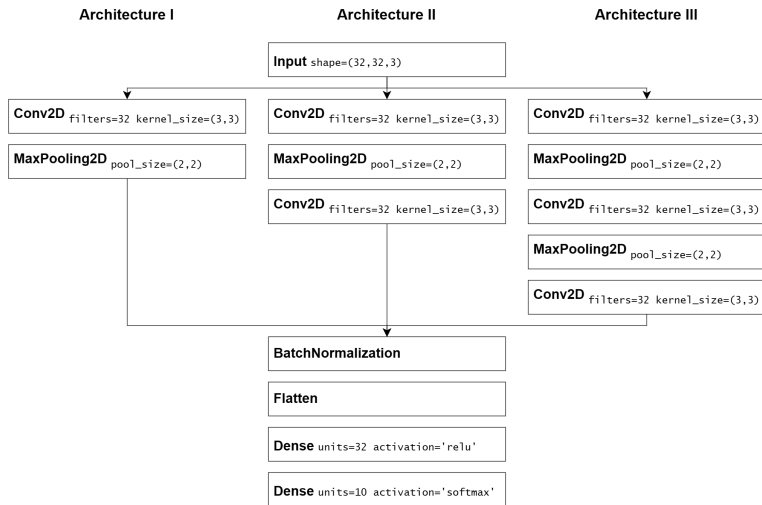


Figure: Three neural network architectures used in the experiments.

Experiments - details

Experiment	Objective	Values
1.1	Optimizer	SGD RMSProp Adam Adadelata
1.2	Learning rate	0.0001 0.0005 0.001 0.005 0.01
1.3	Dropout rate	0.2 0.4 0.6
1.4	Early stopping - patience	2 4 6
2	Augmentation technique	image flips image rotations image shifts mix-up

Table: Details of the experiments.

Experiments - augmentation techniques



Figure: Sample effect of augmentation techniques.

Experiment 0 - architecture

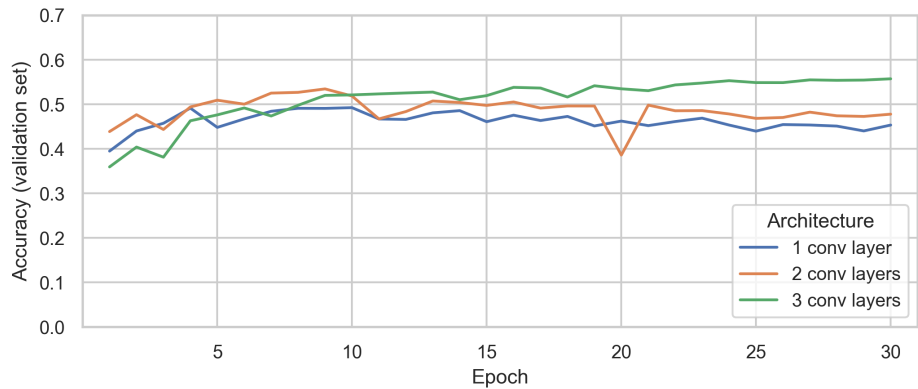


Figure: Accuracy computed on the validation subset for every epoch.

Experiment 1.1 - optimizer




Figure: Accuracy computed on the validation subset for every epoch.

Experiments 1.2 & 1.3 - learning rate and dropout rate

Learning rate	Accuracy	Validation accuracy
0.0001	0.625 (0.007)	0.557 (0.005)
0.0005	0.658 (0.015)	0.574 (0.001)
0.001	0.652 (0.017)	0.568 (0.001)
0.005	0.592 (0.011)	0.542 (0.004)

(a) Experiment 1.2 - learning rate

Drop rate	Accuracy	Validation accuracy
0.2	0.619 (0.020)	0.589 (0.008)
0.4	0.597 (0.007)	0.585 (0.010)
0.6	0.572 (0.008)	0.582 (0.009)

(b) Experiment 1.3 - drop rate

Table: The mean (and standard deviation) of the best values of the accuracy achieved by models.

Experiment 2 - augmentation techniques




Figure: Accuracy computed on the validation subset for every epoch.

Conclusions

- ▶ The project was prepared according to the instructions.
- ▶ The best parameter settings:
 - ▶ three convolutional layers;
 - ▶ Adam optimizer;
 - ▶ learning rate = 0.001;
 - ▶ dropout rate = 0.2;
 - ▶ early stopping patience = 6;
 - ▶ no augmentation techniques.
- ▶ The accuracy achieved on the validation subset (around 0.58) is not so high.
- ▶ It may be worth to construct more complex architecture.