Project II - Transformers

Deep Learning 2024

Kinga Frańczak, 313335
Grzegorz Zakrzewski, 313555

Contents

Description of the research problem

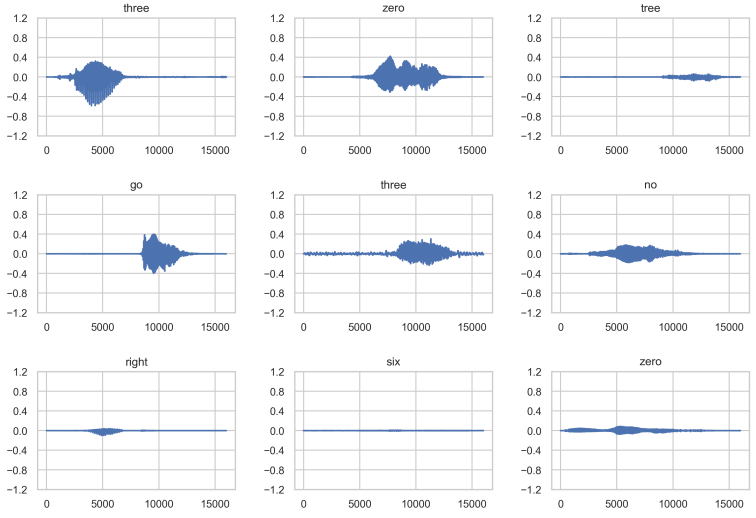

Figure: Sample audio clips from Speech Commands dataset.

Transformer - encoder

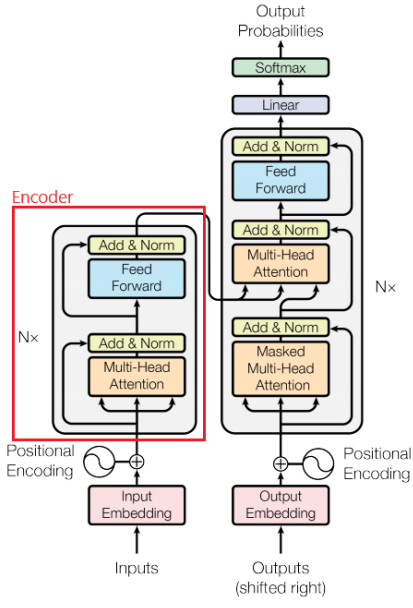

Figure: The original "Attention is All You Need" Transformer diagram.

Experiments - network architectures

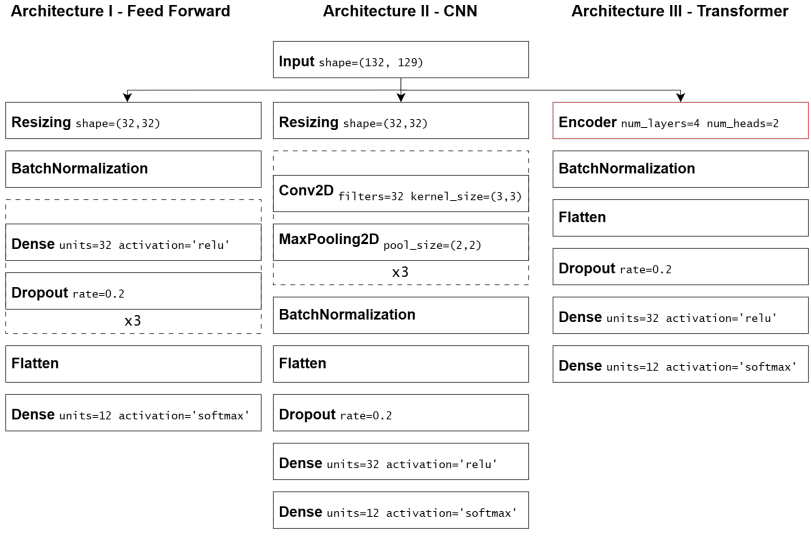

Figure: Three neural network architectures used in the experiments.

Experiments - details

Experiment	Objective	Values
1	Architecture	Simple feed-forward CNN Transformer
2.1	Number of attention heads	2 4 6 8 10
2.2	Number of Encoder sub-layers	2 4 6 8
3	Handling *silence* and *unknown* classes	one network for all classes separate network for special cases

Table: Details of the experiments.

Experiment 1 - architecture

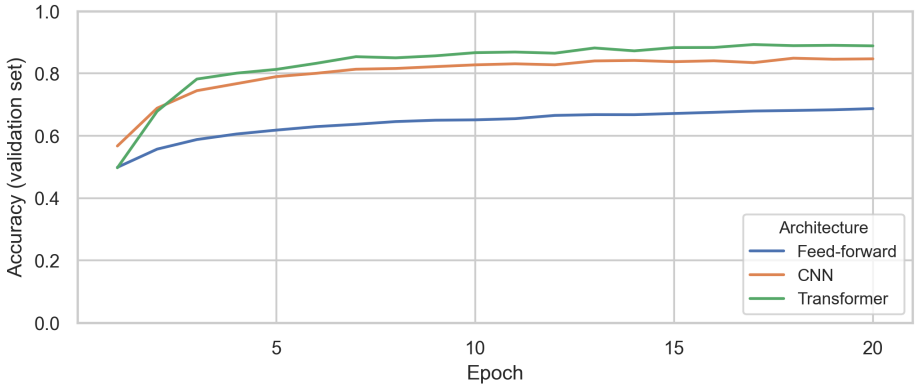

Figure: Accuracy computed on the validation subset for every epoch.

Experiments 2.1 & 2.2 - hyper-parameters

Number of attention heads	Accuracy	Validation accuracy
2	0.930 (0.004)	0.901 (0.003)
4	0.914 (0.008)	0.888 (0.012)
6	0.866 (0.036)	0.845 (0.025)
8	0.794 (0.021)	0.791 (0.016)
10	0.803 (0.023)	0.798 (0.012)

(a) Experiment 2.1 - number of attention heads.

Number of Encoder sub-layers	Accuracy	Validation accuracy
2	0.941 (0.003)	0.875 (0.004)
4	0.930 (0.004)	0.897 (0.006)
6	0.905 (0.015)	0.884 (0.009)
8	0.468 (0.132)	0.424 (0.174)

(b) Experiment 2.2 - number of Encoder sub-layers.

Table: The mean (and standard deviation) of the best values of the accuracy achieved by models.

Experiment 3 - handling *silence* and *unknown* classes.

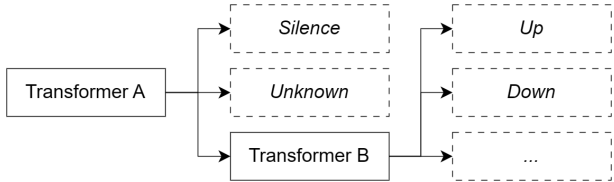

Figure: Experiment 3 - diagram

	Accuracy	Validation accuracy
Transformer A	0.947 (0.015)	0.934 (0.011)
Transformer B	0.940 (0.003)	0.913 (0.005)
Combination	0.944 (0.014)	0.917 (0.011)

Table: Experiment 3 - results.

Confusion matrix

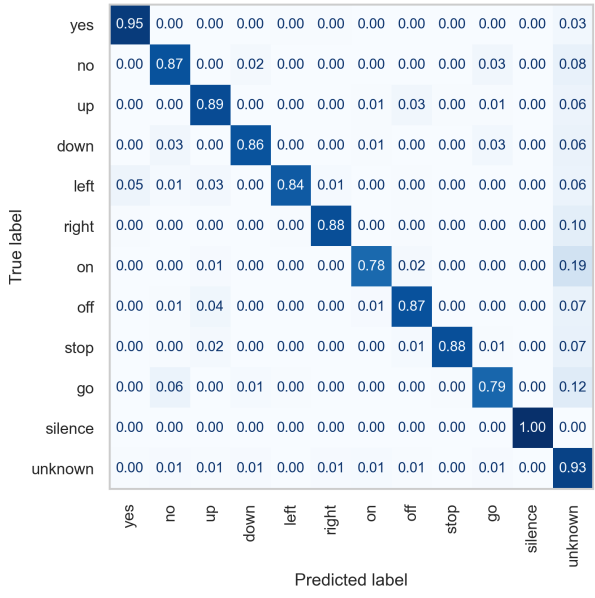

Figure: Confusion matrix prepared on the validation dataset, normalized by true conditions (rows).

Conclusions

- The project was prepared according to the instructions.
- The best parameter settings:
  - number of attention heads = 2;
  - number of Encoder sub-layers = 4;
  - separate network for *silence* and *unknown* classes;
- The accuracy achieved on the validation subset (around 0.91) seems very high.
- The accuracy achieved on Kaggle is 0.67.